

COMP4611: Design and Analysis of
Computer Architectures

Input/Output

Lin Gu

CSE, HKUST

Motivation: Who Cares About I/O?

- CPU Performance: 60% improvement per year
- I/O system performance: < 10% improvement per year
 - sometimes limited by *mechanical* delays (disk I/O)
- 10% IO & 10x CPU => 5x Performance (lose 50%)
10% IO & 100x CPU => 10x Performance (lose 90%)
- I/O bottleneck:
Diminishing value of faster CPUs

Input and Output Devices

I/O devices are diverse with respect to

- Behavior – input, output or storage
- User – human or machine
- Data rate – the rate at which data are transferred between the I/O device and the main memory or processor

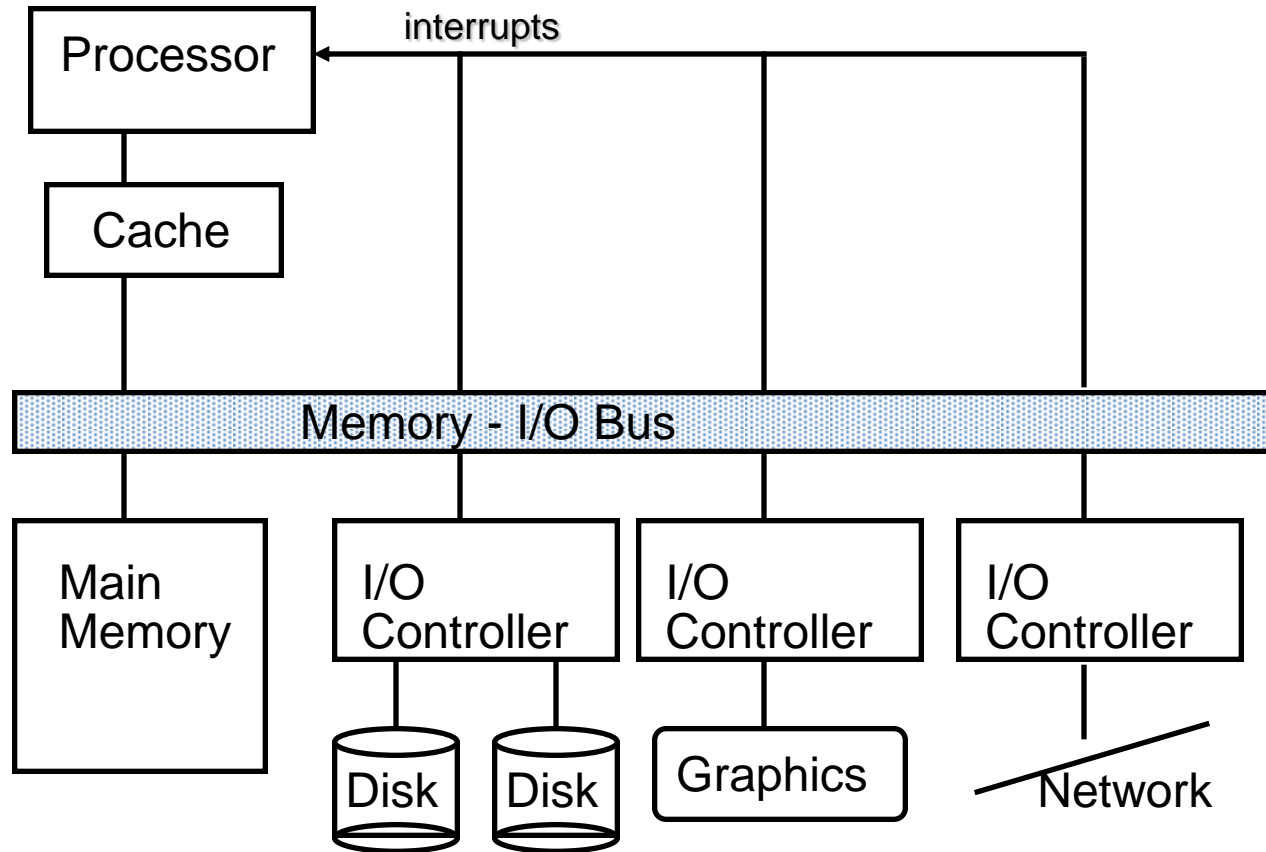
Device	Behavior	User	Data rate (Mb/s)
Keyboard	input	human	0.0001
Mouse	input	human	0.0038
Laser printer	output	human	3.2000
Graphics display	output	human	800.0000-8000.0000
Network/LAN	input or output	machine	10.0000-10000.0000
Magnetic disk	storage	machine	240.0000-2560.0000

9 orders of magnitude
range

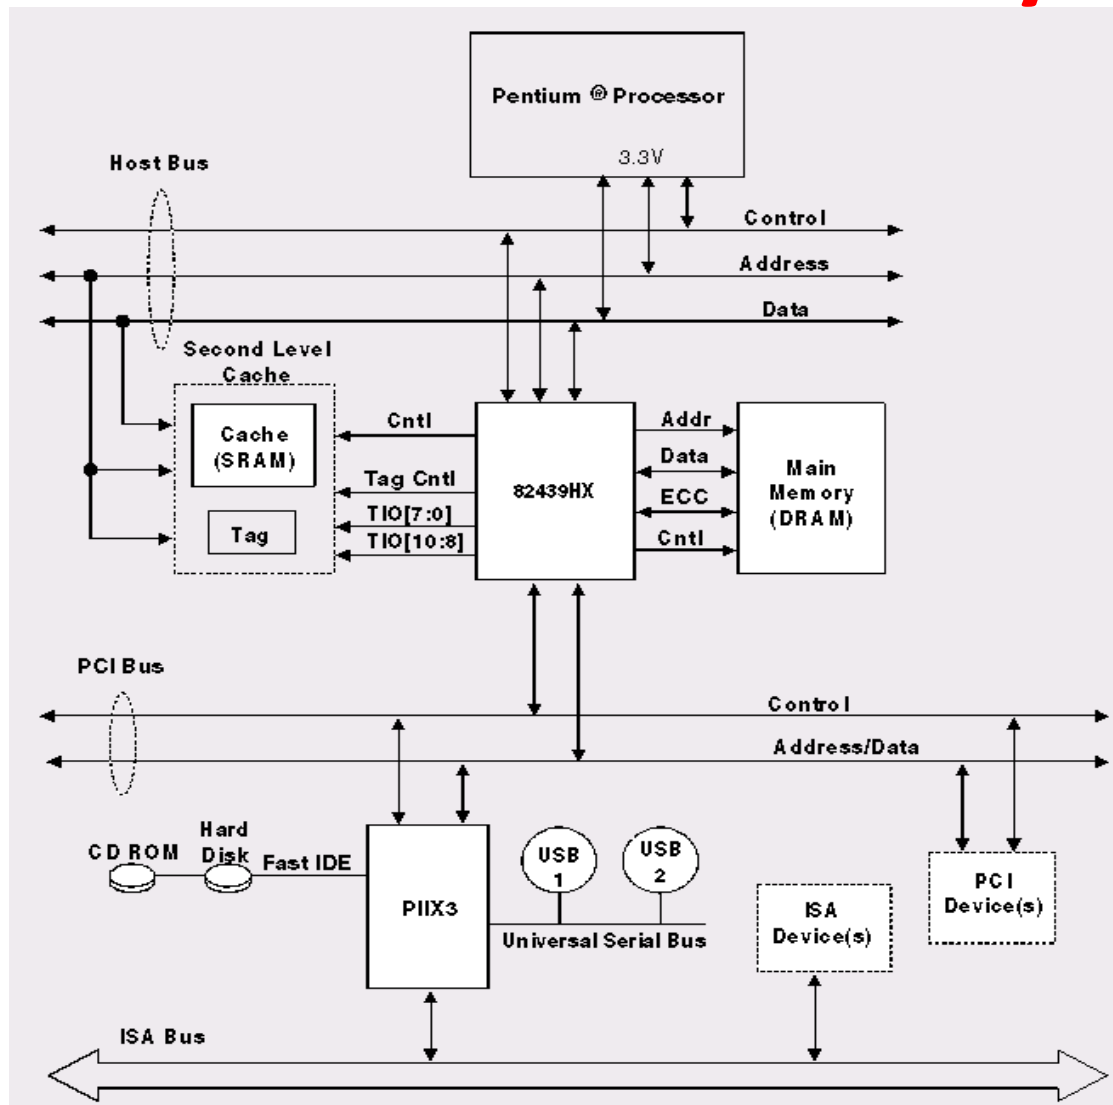
I/O Performance Measures

- **I/O bandwidth** (throughput) – amount of information that can be input (output) and communicated across an interconnect (e.g., a bus) to the processor/memory (I/O device) per unit time
 1. How much data can we move through the system in a certain time?
 2. How many I/O operations can we do per unit time?
- **I/O response time** (latency) – the total elapsed time to accomplish an input or output operation
 - An especially important performance metric in real-time systems
- Many applications require both high throughput and short response times

A Simple System with I/O



A More Recent System



Processor/Memory
Bus


I/O Bus (PCI)

I/O Buses (ISA)


Introduction to I/O

- Most of the I/O devices are *very slow* compared to the cycle time of a CPU.
- The architecture of I/O systems is an active area of R&D.
 - I/O systems can define the performance of a system.
- *Computer architects* strive to design systems that do not tie up the CPU waiting for slow I/O systems (too many applications running simultaneously on processor).

Hard Disk




Hard Disk Performance



- Disk Latency = Seek Time + Rotation Time + Transfer Time + Controller Overhead
- Seek Time depends on the moving distance and moving speed of arms
- Rotation Time depends on how fast the disk rotates and how far sector is from head
- Transfer Time depends on data rate (bandwidth) of disk and size of request


Disk Access Time

I want
block X



block x
in memory

Disk platter



Disk head

Disk arm

Disk access time =

Seek time

+ Rotational delay

+ Transfer time

+ Other delays

Example: Barracuda 180



- 181.6 GB, 3.5 inch disk
- 12 platters, 24 surfaces
- 24,247 cylinders
- 7,200 RPM; (4.2 ms avg. latency)
- 7.4/8.2 ms avg. seek (r/w)
- 65 to 35 MB/s (internal)
- 0.1 ms controller time

Latency =

Queuing Time +
Controller time +
Seek Time +
Rotation Time +
Size / Bandwidth

per access
+
per byte

source: www.seagate.com

Disk Performance Example

Calculate time to read 64 KB (128 sectors) for Barracuda 180 X using advertised performance; sector is on outer track

Disk latency = average seek time + average rotational delay + transfer time + controller overhead

$$= 7.4 \text{ ms} + 0.5 * 1/(7200 \text{ RPM}) \\ + 64 \text{ KB} / (65 \text{ MB/s}) + 0.1 \text{ ms}$$

$$= 7.4 \text{ ms} + 0.5 / (7200 \text{ RPM} / (60000 \text{ ms/M})) \\ + 64 \text{ KB} / (65 \text{ KB/ms}) + 0.1 \text{ ms}$$

$$= 7.4 + 4.2 + 1.0 + 0.1 \text{ ms} = 12.7 \text{ ms}$$

Communication of I/O Devices and Processor

How does the processor command the I/O devices?

– Special I/O instructions


- Must specify both the device (port number) and the command

– For example:

inp `reg, port; register ← port`
out `port, reg; port → register`

Physical
address space

I/O address space



each device gets
one or more
addresses


Communication of I/O Devices and Processor

How does the processor command the I/O devices?

– Memory-mapped I/O

- Portions of the memory address space are assigned to I/O devices
- Read and writes to those memory addresses are interpreted as commands to the I/O devices
- Load/stores to the I/O address space can only be done by the OS

Physical address space



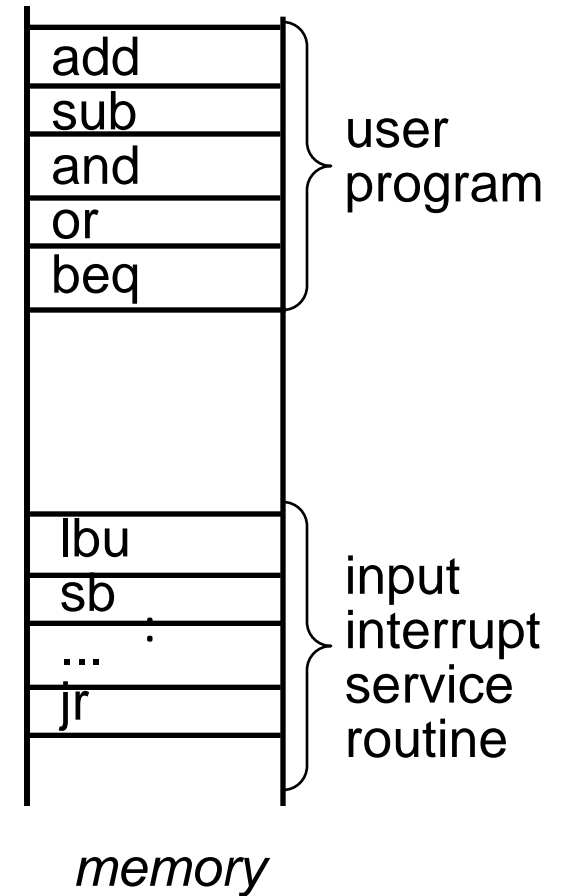

Communication of I/O Devices and Processor

How does the I/O device communicate with the processor?

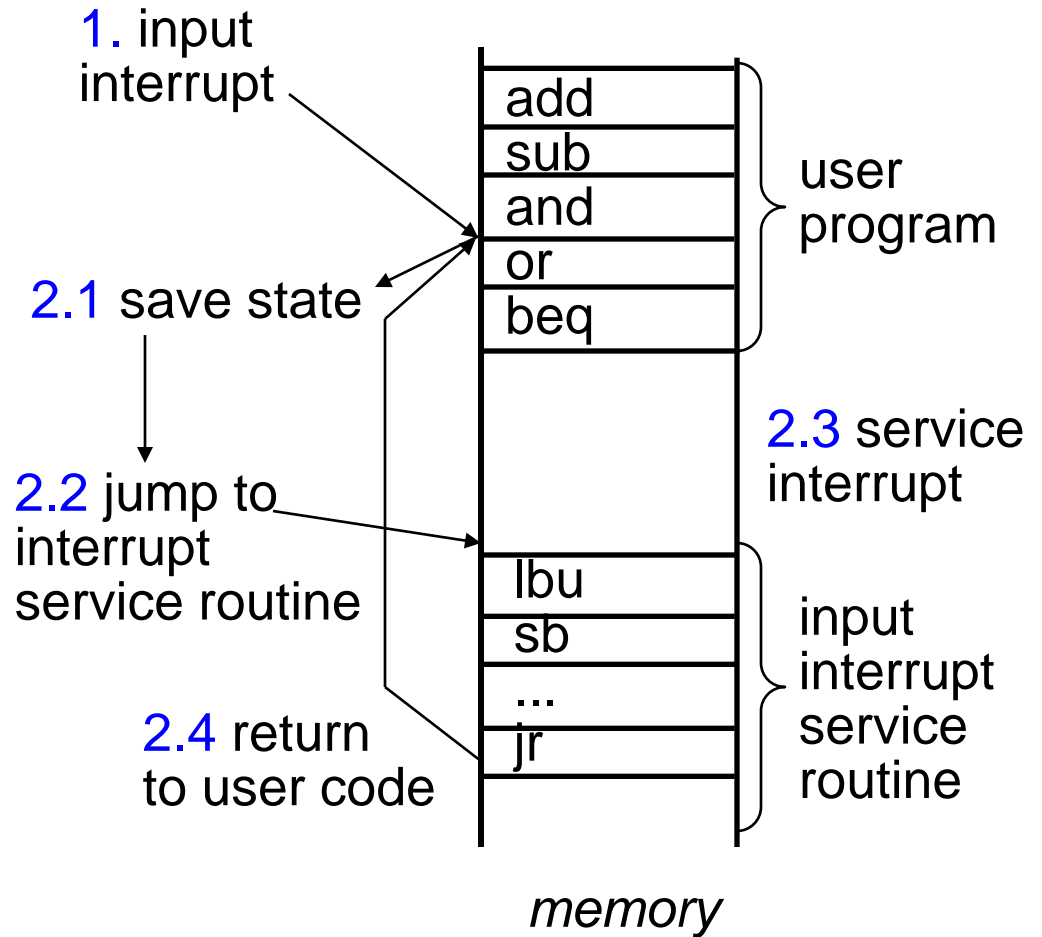

- Polling – the processor periodically checks the status of an I/O device to determine its need for service
 - Processor is totally in control – it also does **all** the work
 - Can waste a lot of processor time due to speed differences
- Interrupt-driven I/O – the I/O device issues an interrupt to the processor to indicate that it needs attention

Interrupt-Driven Input

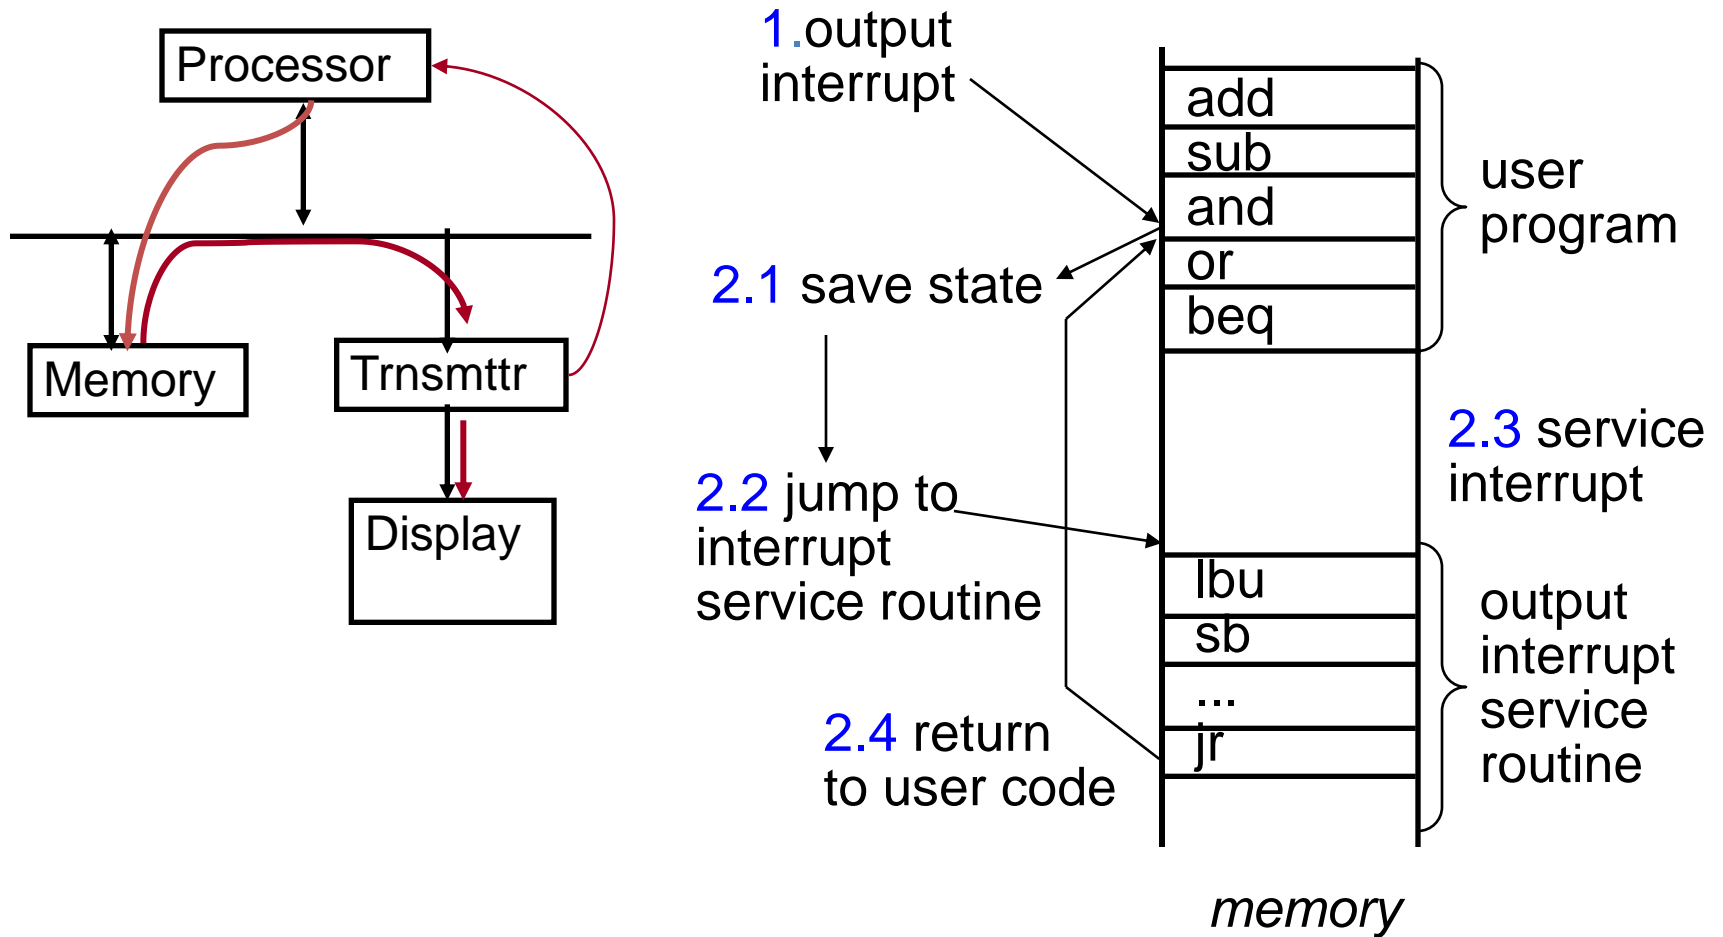
1. input
interrupt



Interrupt-Driven Input



Interrupt-Driven Output




Direct-Memory Access (DMA)

- Interrupt-driven IO relieves the CPU from waiting for every IO event
- But the CPU can still be bogged down if it is used in transferring IO data.
 - Typically **blocks** of bytes.
- For high-bandwidth devices (like disks) interrupt-driven I/O would consume a *lot* of processor cycles

DMA

- DMA – the I/O controller has the ability to transfer data **directly** to/from the memory without involving the processor



DMA

- Consider printing a 60-line by 80-character page
- With no DMA:
 - CPU will be interrupted **4800 times**, once for each character printed.
- With DMA:
 - OS sets up an I/O buffer and CPU writes the characters into the buffer.
 - DMA is commanded (includes the beginning address of the buffer and its size) to print the buffer.
 - DMA takes items from the buffer one-at-a-time and performs everything requested.
 - Once the operation is complete, the DMA sends a **single** interrupt signal to the CPU.

I/O Communication Protocols

- Typically *one* I/O channel controls *multiple* I/O devices.
- We need a two-way communication between the **channel** and the **I/O devices**.
 - The channel needs to send the command/data to the I/O devices.
 - The I/O devices need to send the data/status information to the channel whenever they are ready.

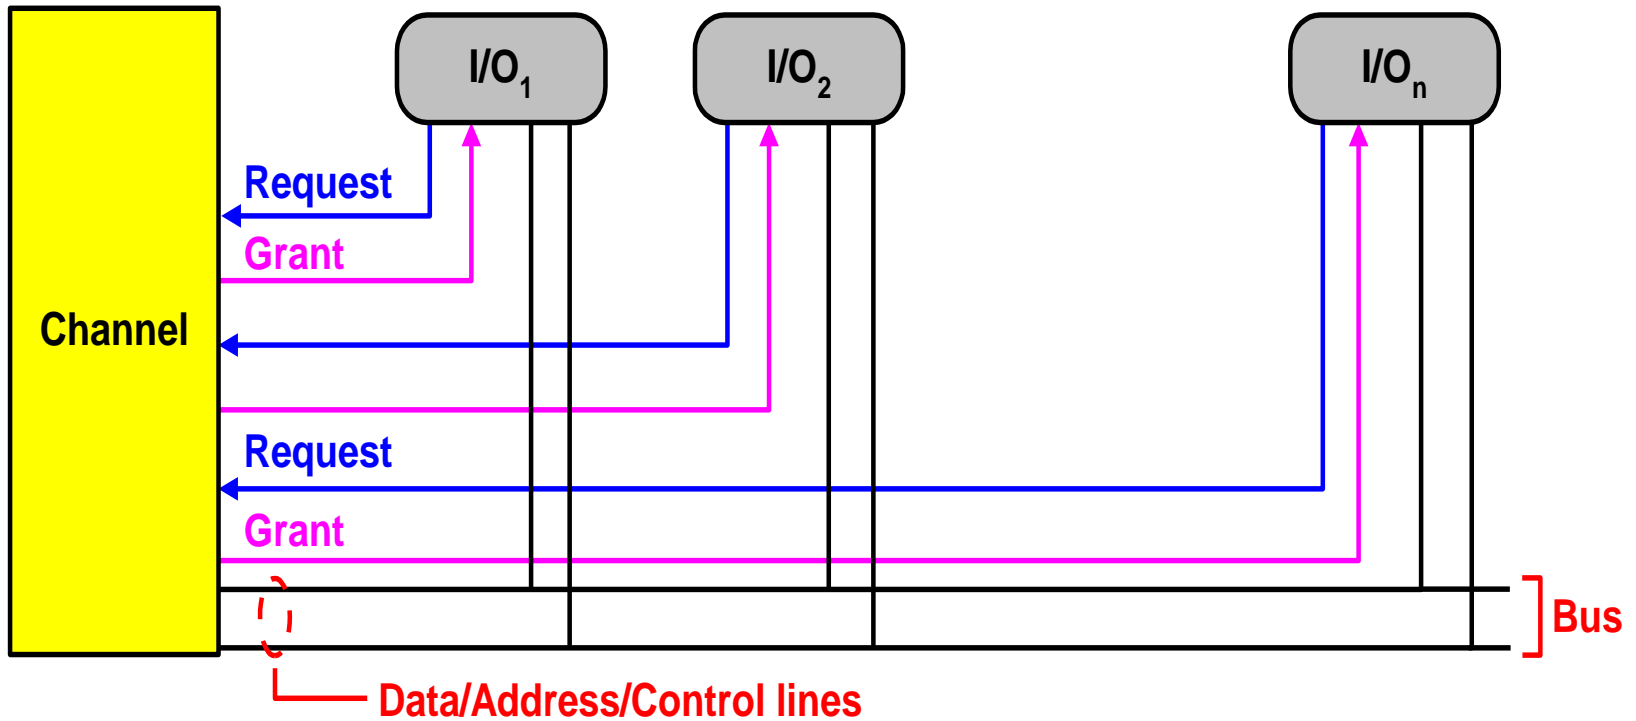
Channel to I/O Device Communication

- Channel sends the address of the device on the bus.
- All devices compare their addresses against this address.
 - Optionally, the device which has matched its address places its own address on the bus again.
 - First, it is an acknowledgement signal to the channel;
 - Second, it is a check of validity of the address.
- The channel then places the I/O command/data on the bus received by the correct I/O device.
- The command/data is queued at the I/O device and is processed whenever the device is ready.

I/O Devices to Channel Communication

- The I/O devices-to-channel communication is more complicated, since now several devices may require *simultaneous* access to the channel.
 - Need arbitration among multiple devices (bus master?)
 - May prefer a priority scheme
- Three methods for providing I/O devices-to-channel communication

Independent Requests



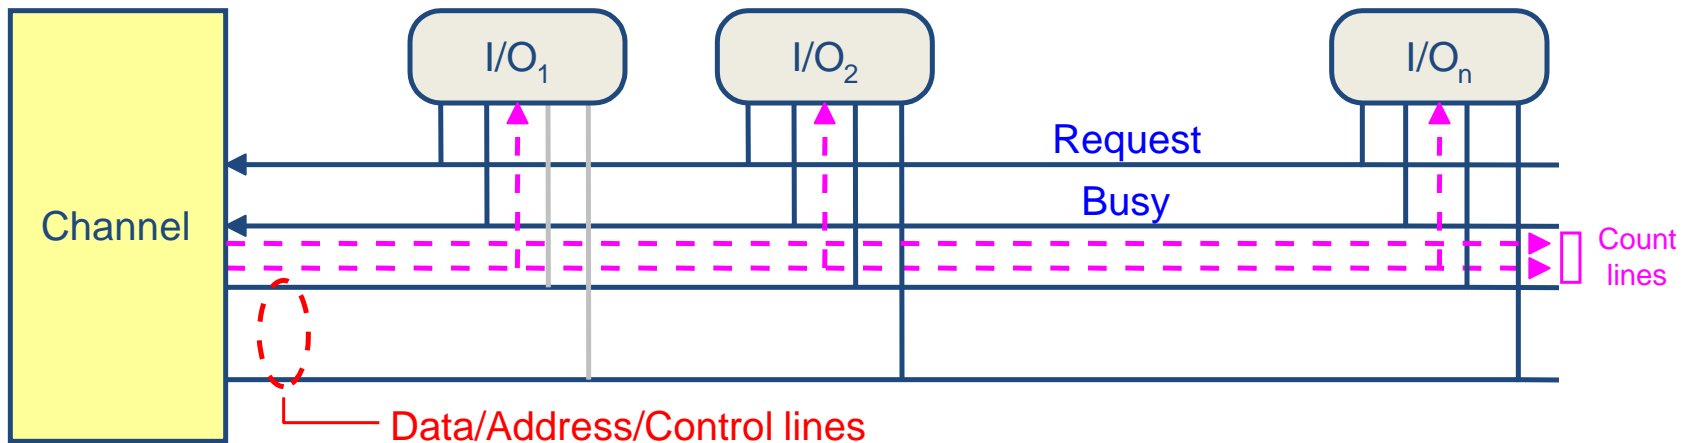
Independent Requests

Each device has its own Request-Grant lines:

- A device sends in its request, the channel responds by granting access
- Only the device that holds the grant signal can access the bus
- When a device finishes access, it lowers its request signal.
- The channel can use either a Priority scheme or Round-Robin scheme to grant the access.

Polling

The channel interrogates (**polls**) the devices to find out which one requested access:

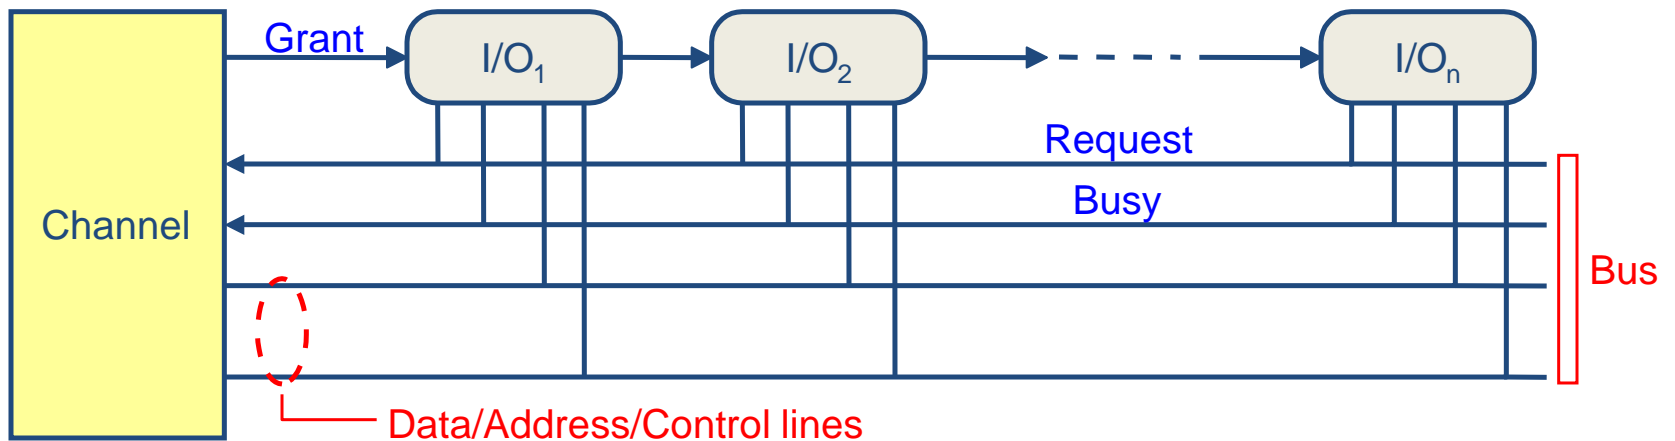


Polling

- Any device requesting access places a signal on request line.
- If the busy signal is off, the channel begins polling the devices to see which one is requesting access.
 - It does this by sequentially sending a count from 1 to n on $\log_2 n$ lines to the devices.
- Whenever a requesting device matches the count against its own number (address), it activates the busy line.
- The channel stops the count (polling) and the device has access over the bus.
- When access is over, the busy line is deactivated and the channel can either continue the count from the last device (**Round-Robin**) or start from the beginning (**priority**).

Daisy Chaining

- Two schemes
- Centralized control (priority scheme)

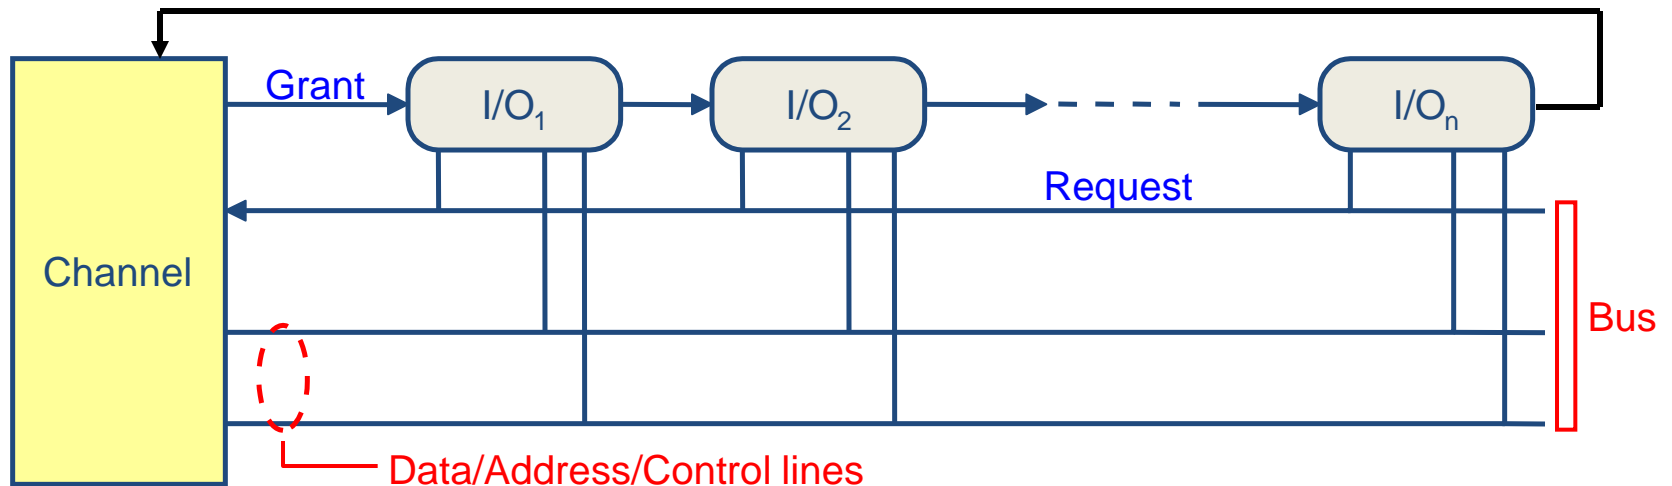


Daisy Chaining

- The I/O devices activate the **request line** for bus access.
- If the bus is not busy (indicated by no signal on busy line), the channel sends a **Grant signal** to the first I/O device (closest to the channel).
 - If the device is not the one that requested the access, it propagates the Grant signal to the next device.
 - If the device is the one that requested an access, it then sends a busy signal on the busy line and begins access to the bus.
- Only a device that holds the Grant signal can access the bus.
- When the device is finished, it resets the busy line.
- The channel honors the requests only if the bus is not busy.
- Obviously, devices **closest** to the channel have a **higher priority** and **block** access requests by lower priority devices.

Daisy Chaining

- Decentralized control (**Round-robin Scheme**)



Daisy Chaining

- The I/O devices send their request.
- The channel activates the Grant line.
- The first I/O device which requested access accepts the Grant signal and has control over the bus.
 - Only the devices that have received the grant signal can have access to the bus.
- When a device is finished with an access, it checks to see if the request line is activated or not.
- If it is activated, the current device sends the Grant signal to the next I/O device (**Round-Robin**) and the process continues.
 - Otherwise, the Grant signal is deactivated.

I/O Buses

- Connect I/O devices (channels) to memory.
 - Many types of devices are connected to a bus.
 - Have a wide range of bandwidth requirements for the devices connected to a bus.
 - Typically follow a bus standard, e.g., PCI, SCSI.
- Clocking schemes:
 - **Synchronous**: The bus includes a clock signal in the control lines and a fixed protocol for address and data relative to the clock
 - **Asynchronous**: The bus is self-timed and uses a handshaking protocol between the sender and receiver

Buses

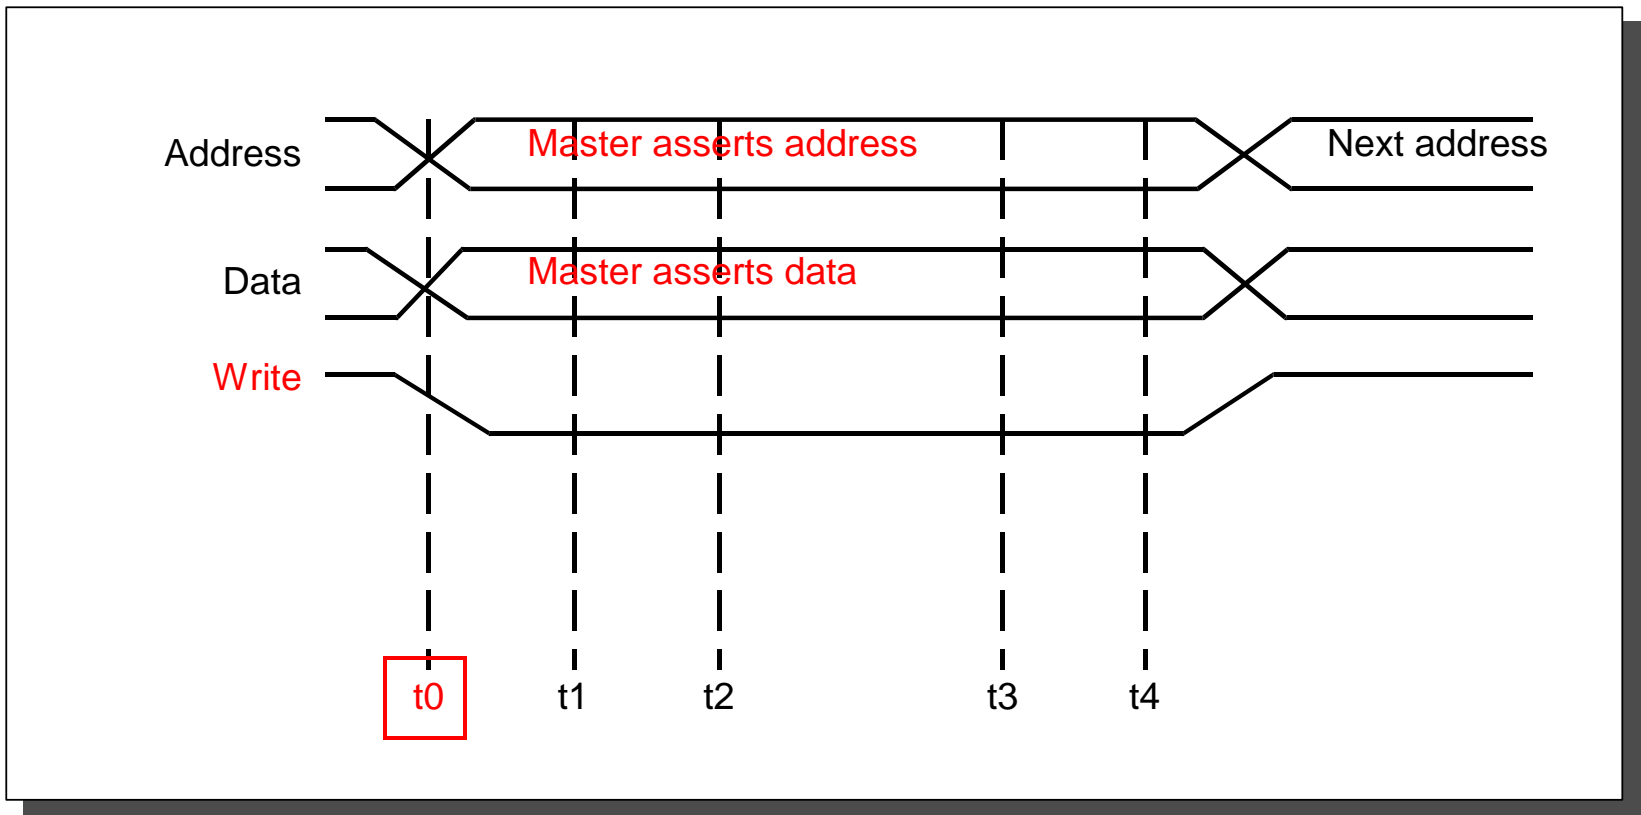
Synchronous buses are fast and inexpensive, but

- All devices on the bus must run at the same clock rate.
- Due to clock-skew problems, buses cannot be long.
- *CPU-Memory* buses are typically implemented as synchronous buses.
 - The front side bus (FSB) clock rate typically determines the clock speed of the memory you must install.

Buses

- **Asynchronous buses** are self-timed and use a handshaking protocol between the sender and receiver.
- This allows the bus to accommodate a wide variety of devices and to lengthen the bus.
- I/O buses are typically asynchronous.
 - A master (e.g., an I/O channel writing into memory) asserts address, data, and control and begins the handshaking process.


I/O Buses



Asynchronous write: master asserts address, data, write buses.

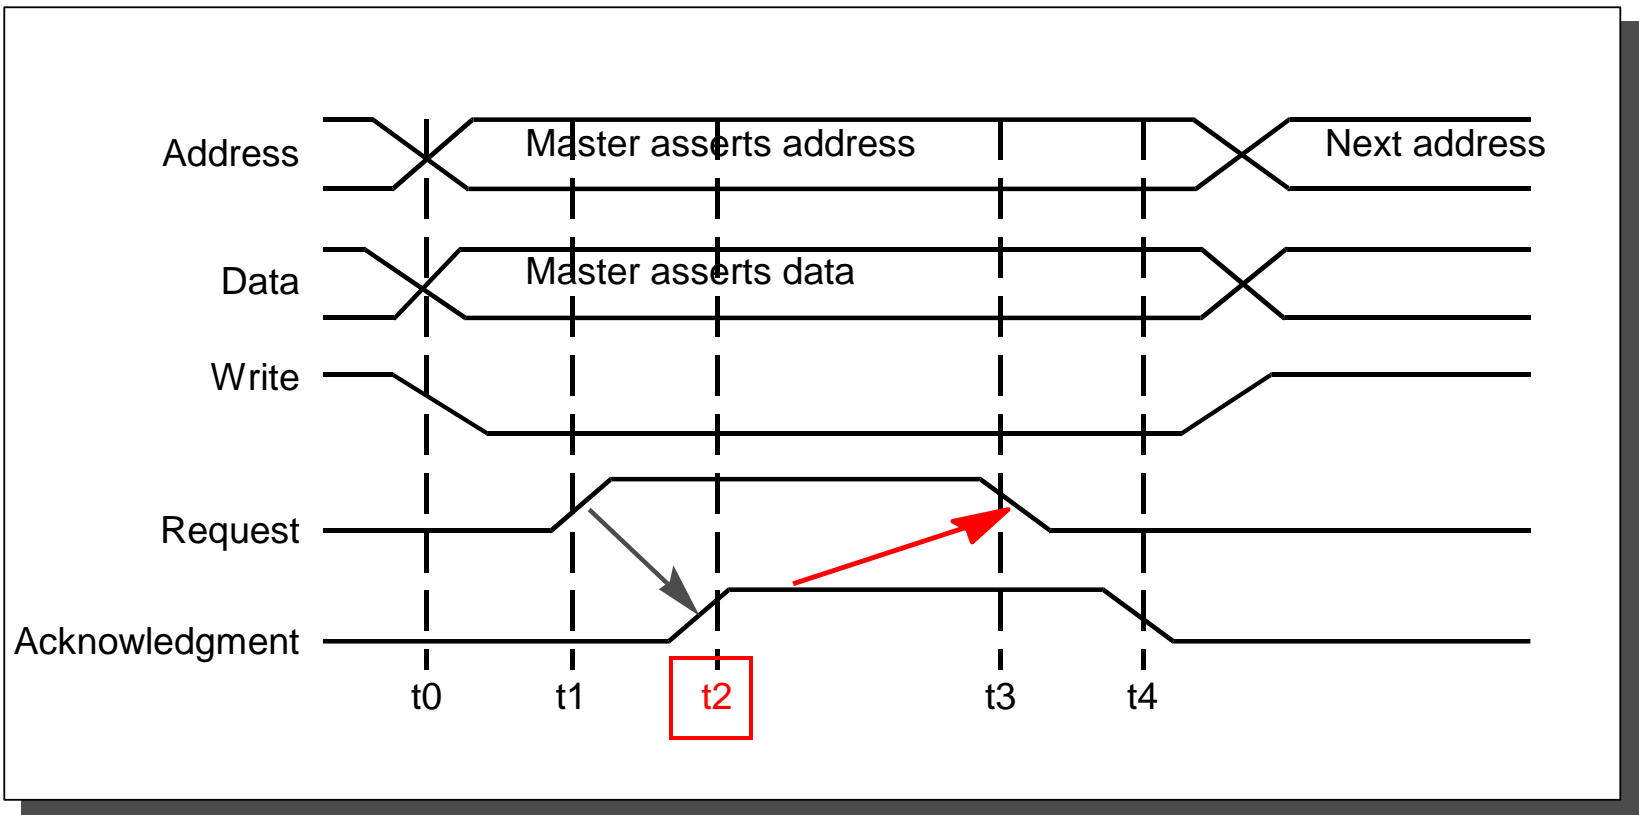
I/O Buses

Asynchronous write: master asserts request, expecting acknowledgement later.




I/O Buses

Asynchronous write: slave (memory) asserts acknowledgment, expecting request to be deasserted later.




I/O Buses

Asynchronous write: master deasserts request and expects the acknowledgement to be deasserted later.



I/O Buses

Asynchronous write: slave deasserts acknowledgement and operation completes.



I/O Bus Examples

Multiple master I/O buses:

	Sun-S bus	IBM MicroChannel	PCI	SCSI 2
Data width	32 bits	32 bits	32 to 64 bits	8 to 16 bits
Clock rate	16 to 25 MHZ	Asynchronous	33 MHZ	10 MHZ or Asynch.
32-bit reads bandwidth	33 MB/Sec	20 MB/Sec	33 MB/Sec	20 MB/sec or 6 MB/Sec
Peak Bandwidth	89 MB/Sec	75 MB/Sec	132 MB/Sec	20 MB/sec or 6 MB/Sec

I/O Bus Examples

Multiple master CPU-memory buses:

	HP Summit	SGI Challenge	Sun XDBus
Data width	128 bits	256 bits	144 bits
Clock rate	60 MHZ	48 MHZ	66 MHZ
Peak Bandwidth	960 MB/Sec	1200 MB/Sec	1056 MB/Sec