

COMP4611: Design and Analysis of
Computer Architectures

Memory System

Virtual Memory

Lin Gu

CSE, HKUST

Why Virtual Memory?

- An example: MIPS64 is a 64-bit architecture allowing an address space defined by 64 bits
 - Maximum address space:
 - $2^{64} = 16 \times 10^{18} = 16,000$ petabytes
 - peta = 10^{15}
- This is several orders of magnitude larger than any realistic and economical (or necessary for that matter) physical memory system

Virtual Memory

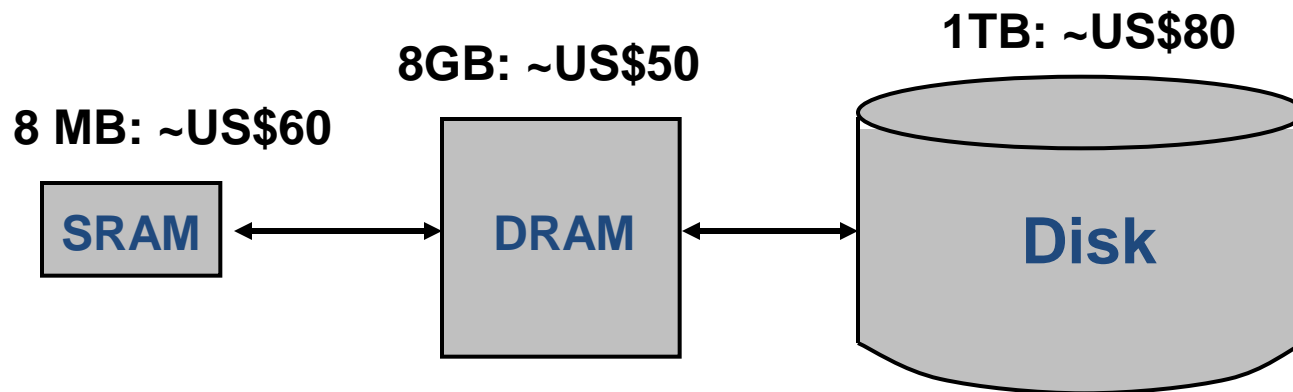
- Originally invented to support program sizes larger than then-available physical memory
 - later on it finds applications in multi-programming and virtual machines
- Virtual memory is as large as the address space allowed by the ISA...but
 - only a portion of the address space resides in physical memory at any given time
 - the rest is kept on disks and brought into physical memory as needed

Motivations for Virtual Memory

- (1) Use Physical DRAM as a Cache for the Disk
 - Address space of a process (program) can exceed physical memory size
 - Only “active” code and data are actually in memory
 - Allocate more memory to process as needed.
 - Sum of address spaces of multiple processes can exceed physical memory
- (2) Simplify Memory Management
 - Multiple processes reside in main memory.
 - Each process with its own address space


Motivation #1: DRAM as “Cache” for Disk

- Full address space is quite large:
 - 32-bit addresses: 0—4,294,967,295 (~ 4 billion bytes)
 - 64-bit addresses: 0—18,446,744,073,709,551,615 (~ 16,000 petabytes)
- Disk storage is much cheaper than DRAM storage
 - 1TB DRAM: ~ US\$6,400
 - 1TB disk: ~ US\$80
- To access large amounts of data in a cost-effective manner, the bulk of the data must be stored on disk




DRAM vs. SRAM as a “Cache”

- DRAM vs. disk is more extreme than SRAM vs. DRAM
 - Access latencies:
 - DRAM ~10X slower than SRAM
 - Disk ~**100,000X** slower than DRAM
 - Design decisions made for DRAM caches driven by enormous cost for misses




Multi-programming



- Multiple programs run concurrently, ideally each with the illusion that it owns all the machine resources
- Each program requires memory space in the main memory

Multiprogramming and Memory Protection

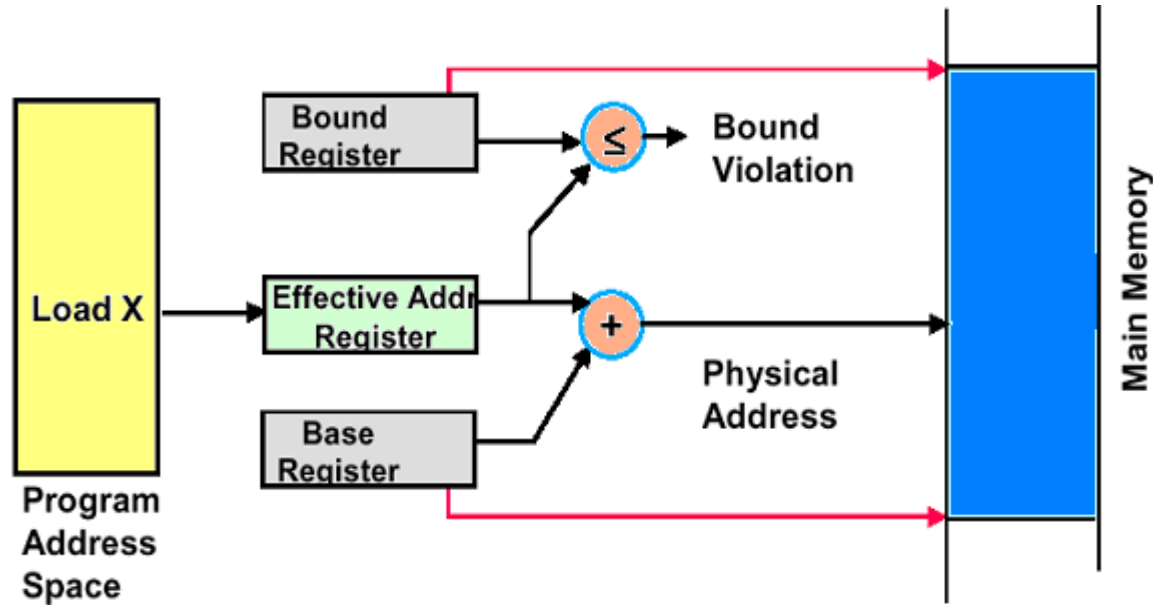


- In the early machines, I/O operations were slow
 - having the CPU wait pending completion of I/O operation was a waste of precious machine cycles
- Higher throughput is achieved if CPU and I/O of 2 or more programs are overlapped
- How to overlap execution of multiple programs?
 - use multiprogramming
- How to protect programs from one another in a multiprogramming environment?
 - use a bound register?

Away With Absolute Addressing

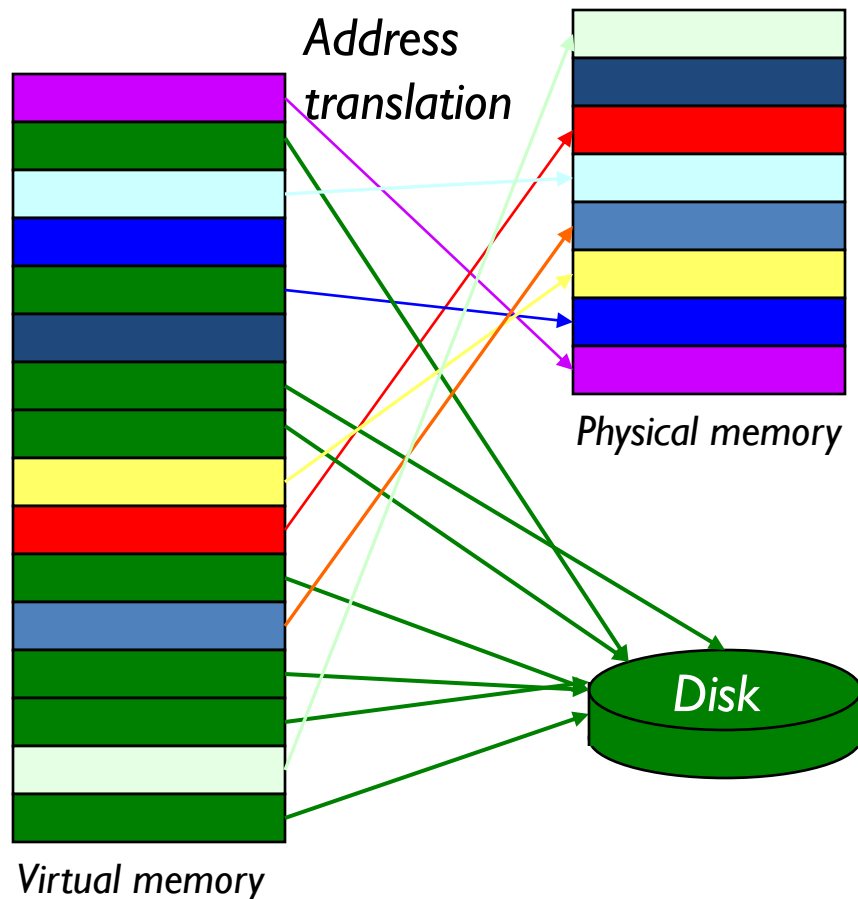
- In early (1950) machines only one program ran at any given time, with unrestricted access to the entire machine (RAM + I/O devices)
- Addresses in a program were location-dependent
 - they depended upon where the program was to be loaded in memory
 - so when programmers wrote code, they had to have a pretty good idea where in memory they should load the program
 - but it was more convenient for programmers to write location-independent programs
- **How to achieve location-independence?**
 - a base register was used to support re-locatable code

Base & Bound Registers



- Base & bound registers loaded from the process table when the program is context-switched
- Permits multiple memory-resident concurrent programs
- Protects programs from one another

Virtual Memory



- Address translation
- Program “sees” entire VM address space
- Program is run in physical memory which is typically smaller than the address space
- Pages of address space are swapped in/out of disk storage as needed
- Strictly speaking VM is required to overcome limitation on the size of physical memory
 - but VM is extended in a natural way to support multiprogramming & memory protection
 - There are machines where the physical memory space is larger than the virtual memory space (e.g., PDP-11)

Advantages of Virtual Memory

- Abstraction of large and flat memory
 - program has a consistent view of a contiguous memory, even though physical memory is scrambled
 - Allows multiprogramming
 - relocation: allows the same program to run in any location in physical memory
- Protection
 - different processes are protected from each other
 - different pages can have different behavior (read-only; user/supervisor)
 - kernel code/data protected from user programs
- Sharing
 - can map same physical memory to multiple processes (shared memory)

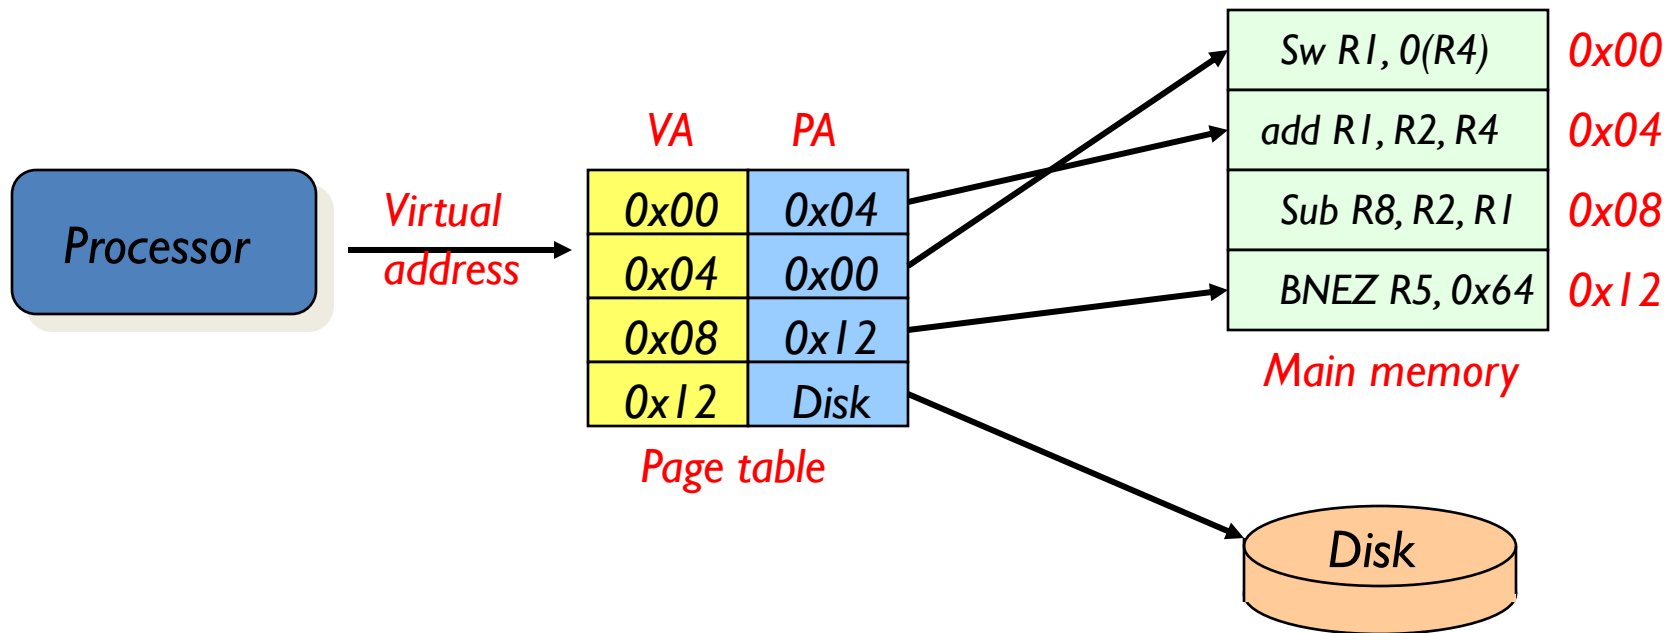
How VM Works

- On program startup
 - OS loads part of the program into RAM; this includes enough code to start execution
 - if program size exceeds allocated RAM space the remainder is maintained on disk
- During execution
 - if program needs a code/data not resident in RAM, it fetches the segment from disk into RAM
 - if there is not enough room in RAM, some resident code/data is evicted from memory to make room
 - if evicted contents are “dirty”, they are written to disk

Address Translation

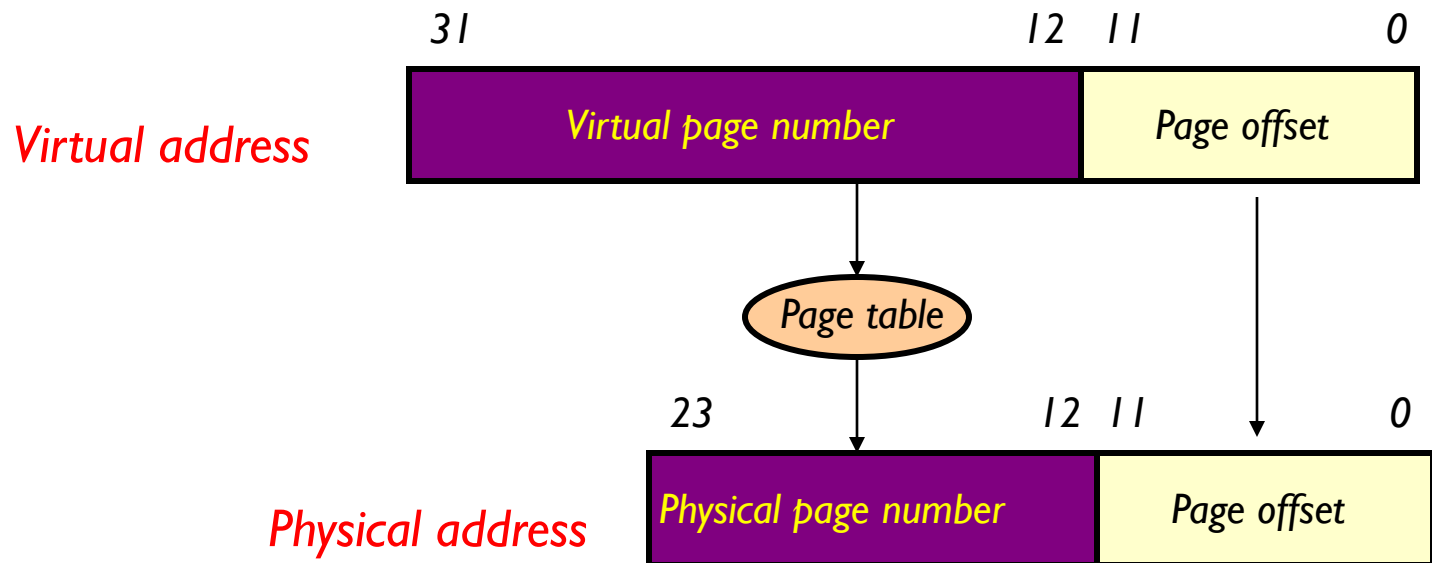
VA-to-PA Address Translation

- Programs use virtual addresses (VA) for data & instructions
- VA translated to physical addresses (PA)
 - Usually organized in pages (paging) or segments
 - Paging systems use a “page table” for the translation
- Instruction/data fetched/updated in memory




Page Table

- Memory organized in pages (similar to blocks in cache)
- Page size is usually 4-8 Kbytes



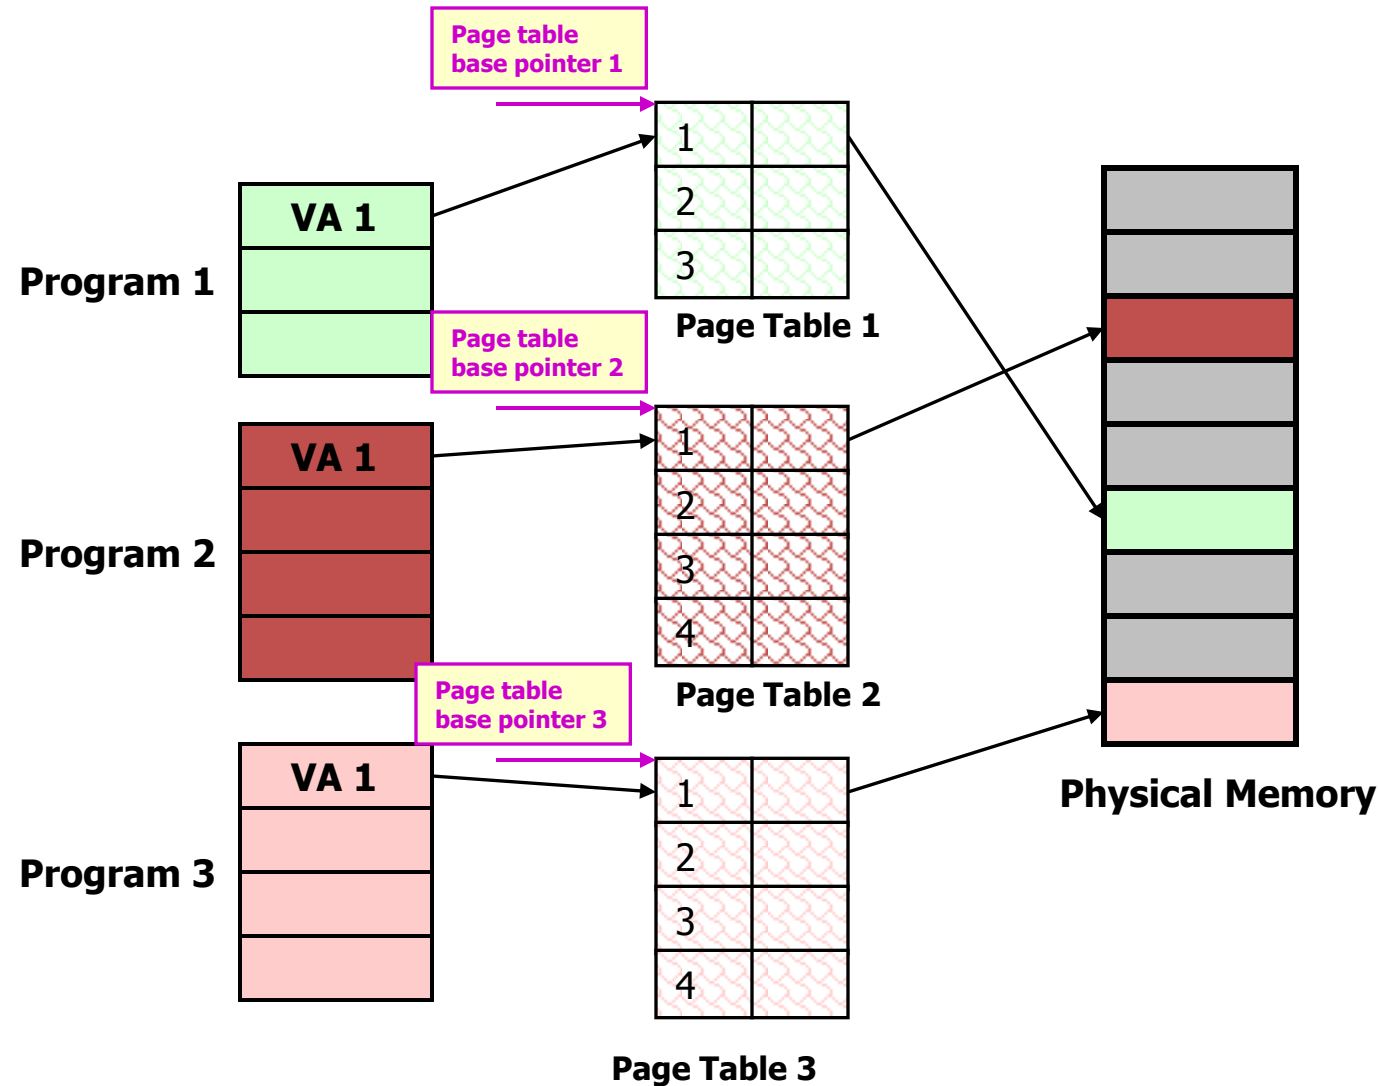
VA to PA Translation



- Page table maps the base address of virtual and physical page frames
- Page offset need not be used in VA=>PA translation because virtual and physical block sizes are the same



Multiprogramming View of VM

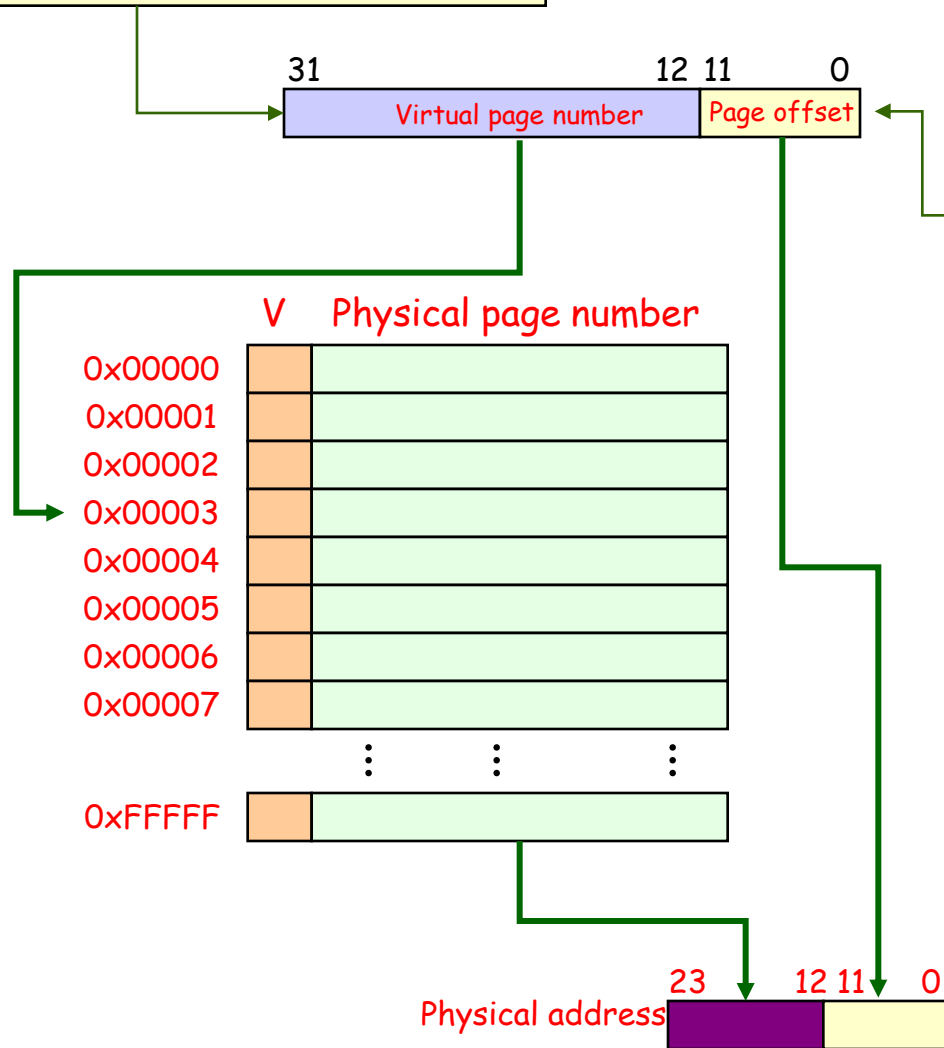


- Each application has a separate page table
- Page table contains an entry for each page
- Page tables are identified by the base pointers maintained by OS for individual processes

Page Table Structure

Virtual page number \Leftrightarrow page table size
20 bits: ~1 million

Page offset \Leftrightarrow page size
12 bits: 4096 bytes



Determining Page Table Size


- Assume
 - 32-bit virtual address
 - 30-bit physical address
 - 4 KB pages => 12 bit page offset
 - Each page table entry is one word (4 bytes)
- How large is the page table?
 - Virtual page number = $32 - 12 = 20$ bits
 - Number of entries = number of pages = 2^{20}
 - Total size = number of entries x bytes/entry
 $= 2^{20} \times 4 = 4 \text{ Mbytes}$
 - Each process running needs its own page table

Page Fault


- How is it known whether the page is in memory?
 - Maintain a valid bit per page table entry
 - valid bit is set to **INVALID** if the page is not in memory
 - valid bit is set to **VALID** if the page is in memory
- Page fault occurs when a page is not in memory
 - fault results in OS fetching the page from disk into DRAM
 - if DRAM is full, OS must evict a page (**victim**) to make room
 - if victim is **dirty** OS updates the page on disk before fetch
 - OS changes page table to reflect turnover
- After a page fault and page-fetching, execution resumes at the instruction which caused the fault

Page Faults (like “Cache Misses”)

Before fault

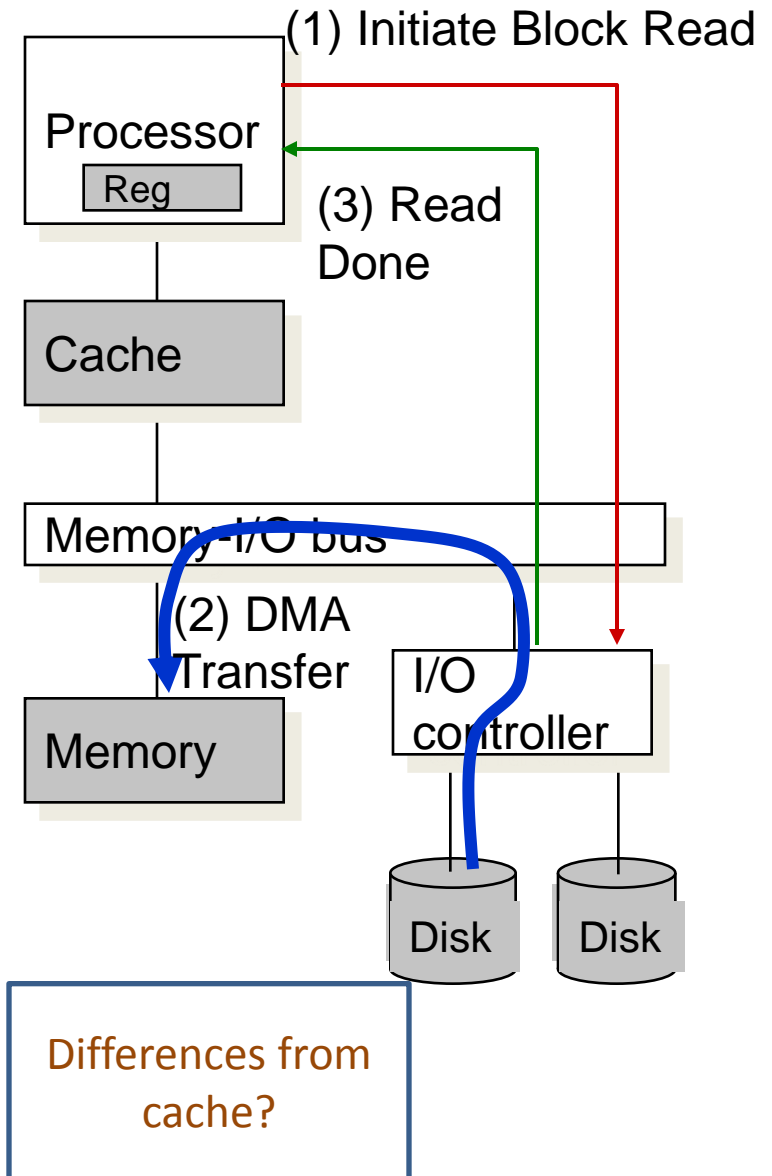


After fault

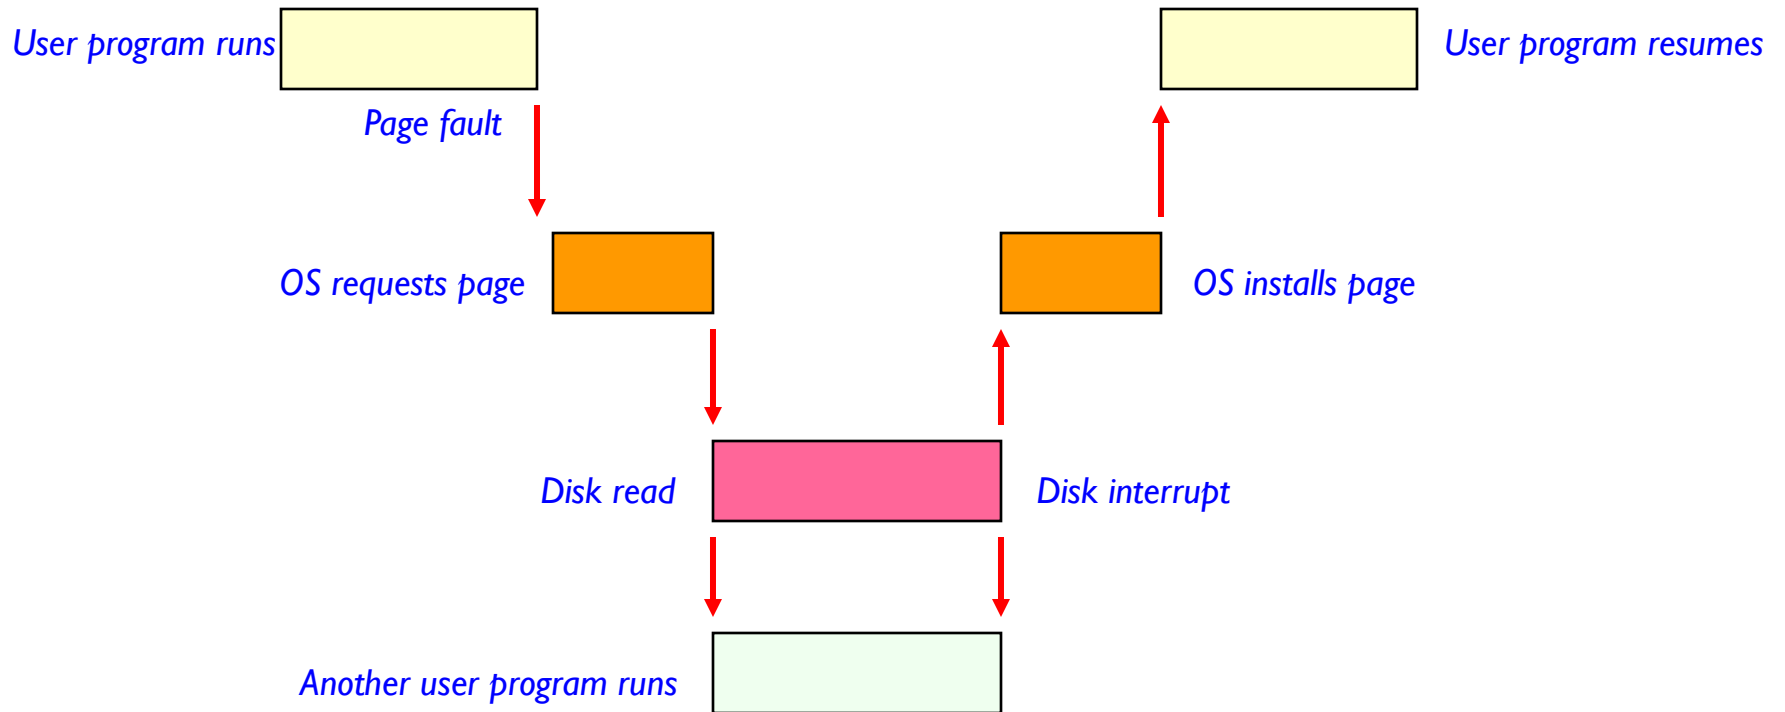


Page Fault Handling

- **Processor signals Controller**
 - Read block of length P starting at disk address X and store data starting at memory address Y
- **Read**
 - Direct Memory Access (DMA)
 - Under control of I/O controller
- **I / O controller signals completion**
 - Interrupts processor
 - OS resumes suspended process



Page Fault Handling

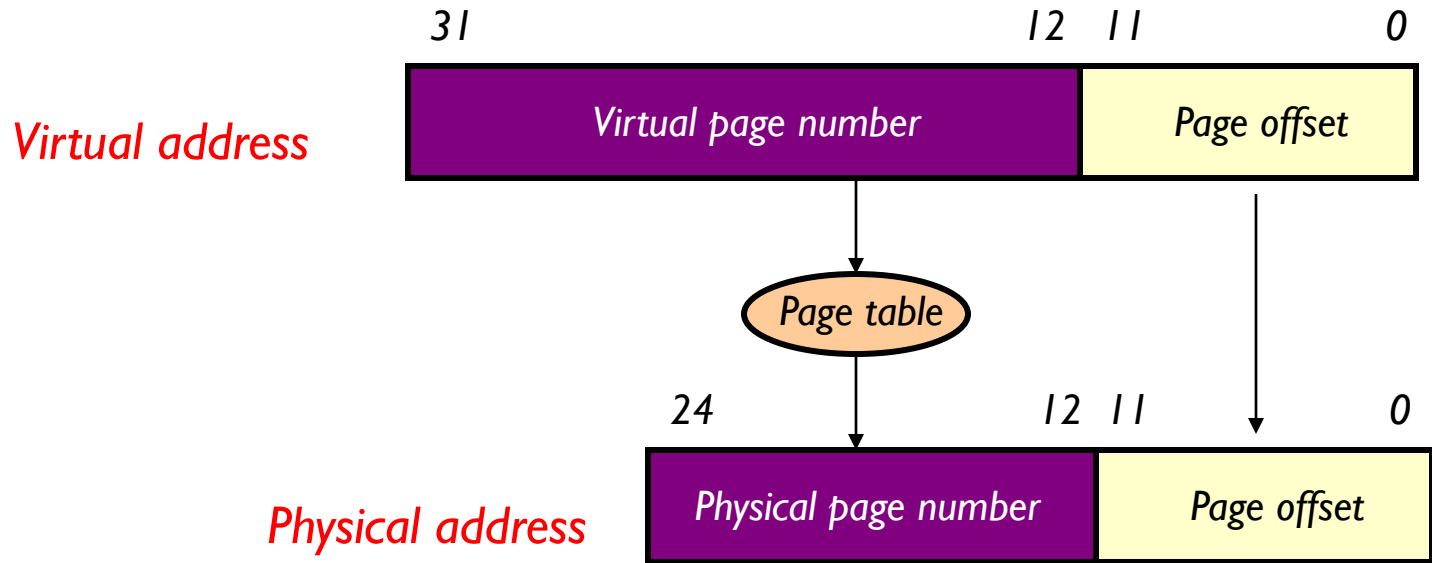


Accelerating Virtual Memory Operations

Virtual Memory Hardware

- Protection via virtual memory
 - Keeps processes in their own memory space
- Role of architecture:
 - Provide user mode and supervisor mode
 - Protect certain aspects of CPU state
 - Provide mechanisms for switching between user mode and supervisor mode
 - Provide mechanisms to limit memory accesses
 - Provide TLB to accelerate the address translation


VM => PM Translation



- Each program memory reference requires two memory accesses: one for VM => PM mapping and one for actual data/instruction
 - must make page table lookup as fast as possible
- Page table too big to keep in fast memory (SRAM) in its entirety
 - store page table in main memory
 - cache a portion of the page table in TLB (Translation Look-Aside Buffer)

Virtual Addressing with a Cache

- Thus it takes an *extra* memory access to translate a VA to a PA

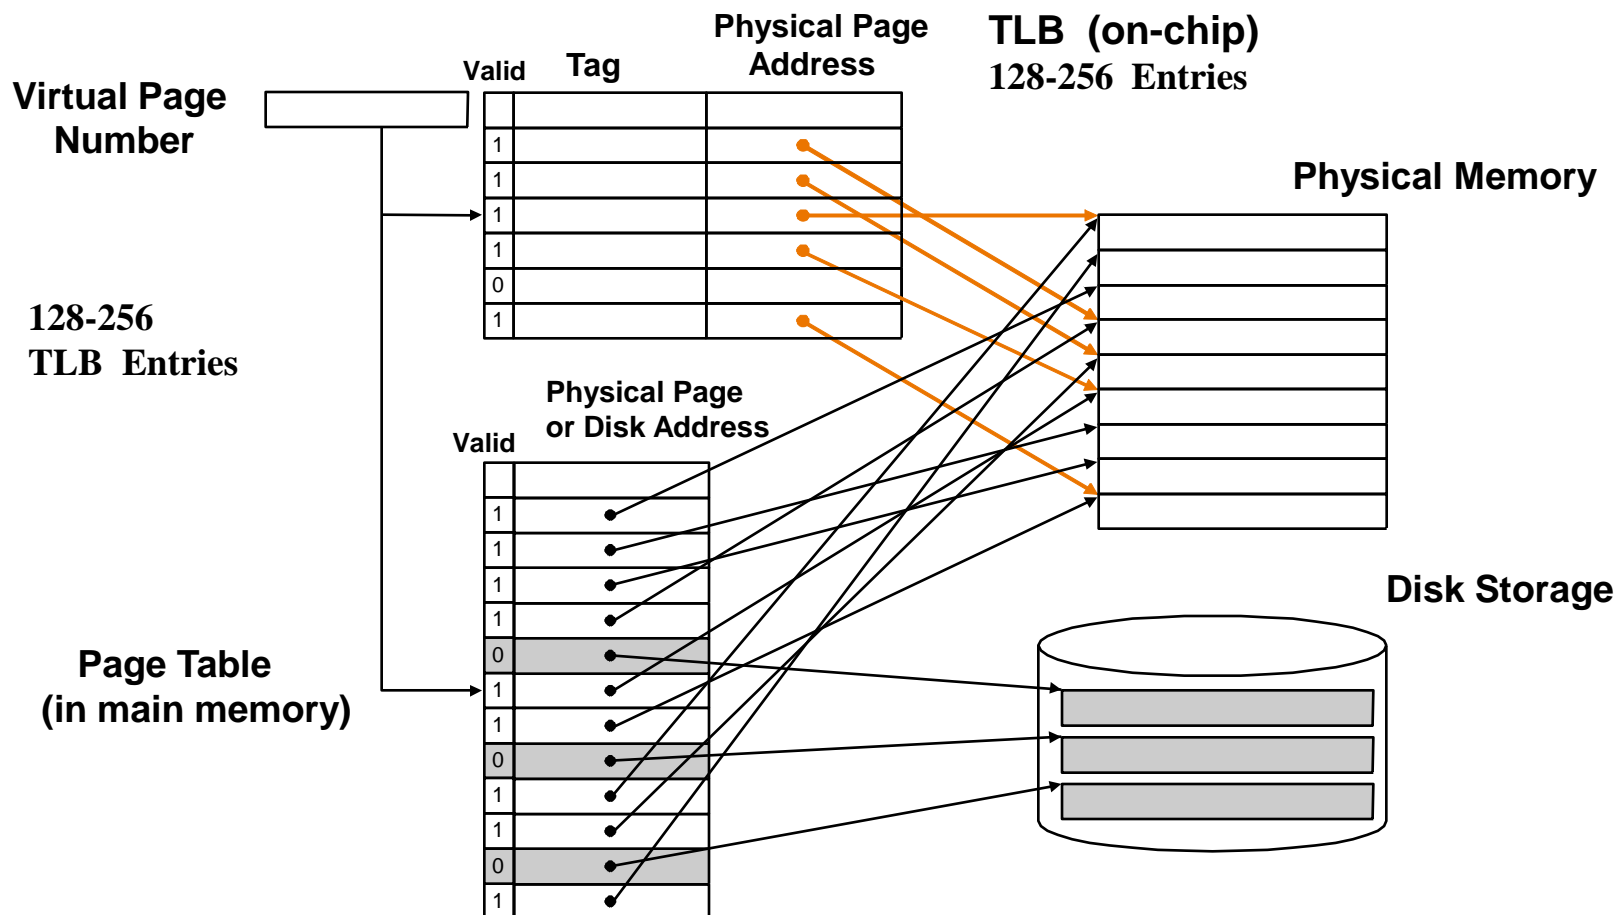


- This makes memory (cache) accesses **very expensive** (if every access was really *two* accesses)
- Translation Lookaside Buffer (TLB) keeps track of recently used address mappings to avoid having to do a page table lookup

Speeding Up Address Translation:


Translation Lookaside Buffer (TLB)

- TLB: A small and fast on-chip memory structure used for address translations.
- If a virtual address is found in TLB (a TLB hit), the page table in main memory is not accessed.

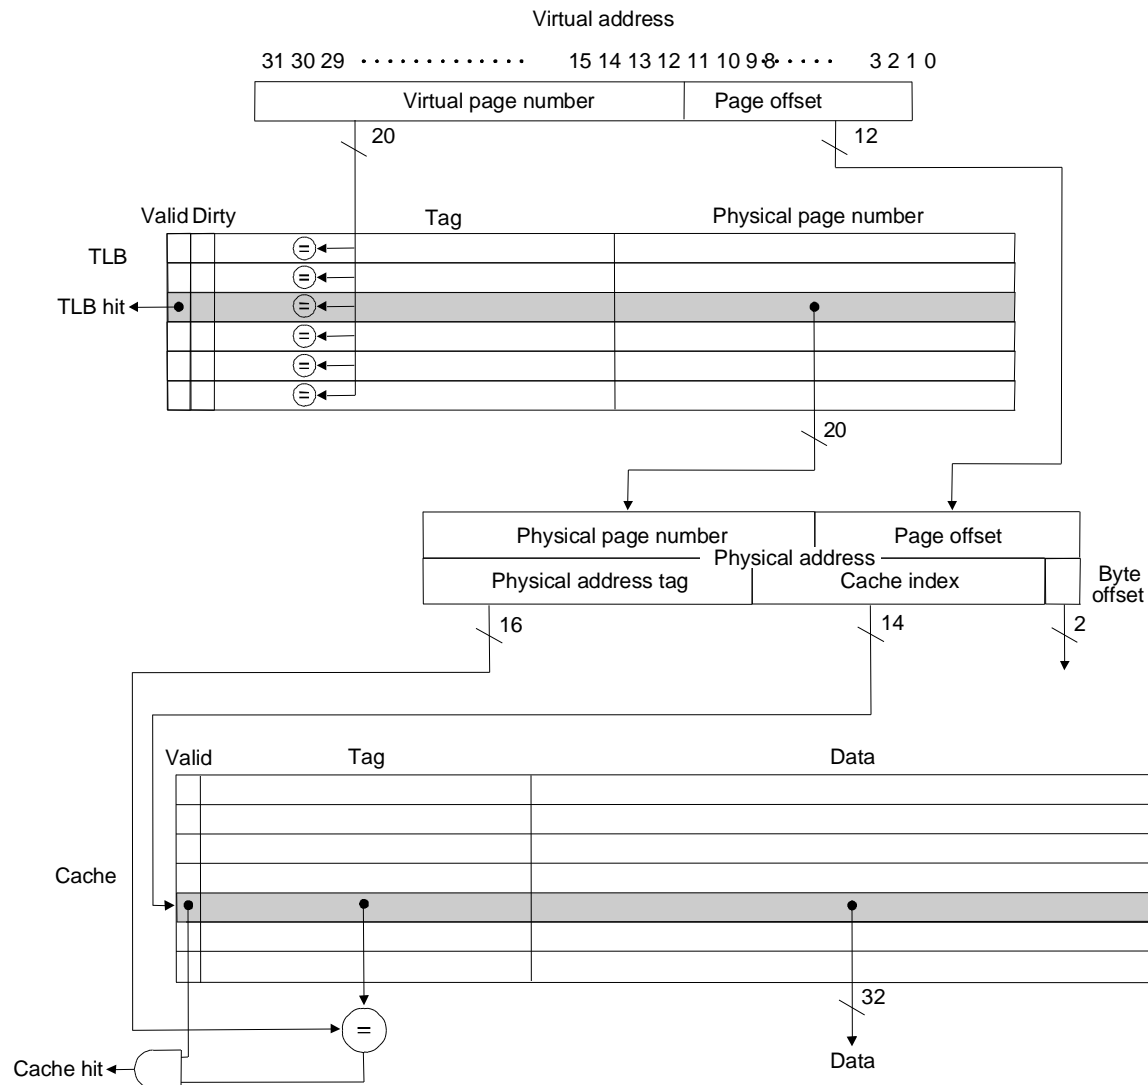


Translation Look-Aside Table (TLB)

- TLB maintains a list of most-recently used pages
- Similar to instruction & data cache
 - virtual address is used to index into the TLB
 - TLB entry contains physical address
 - takes ~ 1 clock cycle
- What if VM=>PM lookup and data/instruction lookup takes more than 1 clock cycle?
 - To avoid TLB lookup, remember the last VM => PM translation
 - if same page is referenced, TLB lookup can be avoided




TLB / Cache Interaction



Techniques for Fast Address Translation

- Just like any other cache, the TLB can be organized as fully associative, set associative, or direct mapped
- TLBs are usually small, typically not more than 128 - 256 entries even on high end machines. This permits fully associative lookup on these machines. Most mid-range machines use small n-way set associative organizations.




TLB and Context Switch

- In multi-programming we need to use TLB for the active process
 - What to do with TLB at context switch?
- Too costly to clear TLB on every context switch
- Keep track of PTE in TLB per process using ASID

TLB Organization

Tag



<i>Virtual address</i>	<i>Physical address</i>	<i>Used</i>	<i>Dirty</i>	<i>Valid</i>	<i>Access</i>	<i>ASID</i>
0xFA00	0x0003	N	Y	Y	R/W	34
0x0040	0x0010	Y	N	Y	R	0
0x0041	0x0011	Y	N	Y	R	0

Additional info per page

- Dirty: page modified (if the page is swapped out, should the disk copy be updated?)
- Valid: TLB entry valid
- Used: Recently used or not (for selecting a page for eviction)
- Access: Read/Write
- ASID: Address Space ID

Example

Consider a virtual memory system with the following properties:

- 32-bit virtual address (4 Gbytes)
 - 26-bit physical memory (64 Mbytes)
 - 4 Kbyte pages (12 bits)
- a. How big is the page table for this memory system, assuming each page table entry has 4 overhead bits (*valid*, *protection*, *dirty*, *use*) and holds one physical page number? Give your answer in bits.
- b. Approximately how many pages of user data, at maximum, can we have in physical memory?
- c. If we wanted to keep the page table under 1Mbit in size, how big would we have to make the pages? Assume a page needs to be 2^K bytes, and find the correct value of K.

Part a

$$AddressSpaceSize = 2^{32} = 4 \times 2^{30} \approx 4 \times 10^9 \text{ Bytes}$$

$$NumberOfPages_{virtual} = \frac{AddressSpaceSize}{PageSize} = \frac{2^{32}}{4 \times 1000} \approx \frac{2^{32}}{2^{12}} = 2^{20}$$

$$NumberOfPages_{virtual} \approx 1,000,000$$

The page table must have 1,000,000 entries. To compute the size in bits note that the page offset need not be stored in the page table. There are 4 Kbytes per page requiring 12 bits of offset ($4000 = 4 \times 10^3 \approx 2^{12}$). Thus each page table entry must map 20 bits of virtual page number ($32 - 12 = 20$) to ($26 - 12 = 14$) bits of physical page number. Taking into account the 4 overhead bits, each entry requires 18 bits. With 1,000,000 entries, the page table size is 18 megabits.

Part b

$$\textit{NumberOfPages}_{\textit{physical}} = \frac{\textit{AddressSpaceSize}_{\textit{physical}}}{\textit{PageSize}} = \frac{2^{26}}{4 \times 1000} \approx \frac{2^{26}}{2^{12}} = 2^{14}$$

$$\textit{NumberOfPages}_{\textit{physical}} \approx 16,000$$

Part c

$$PageTableSize = 1 \text{ Megabits} = \frac{10^6}{8} \text{ Megabytes} \approx \frac{2^{20}}{2^3} = 2^{17} \text{ Megabytes}$$

Assuming that the page is 2^K bytes, we need K bits for page offset in both the physical and virtual addresses. Thus each page entry requires:

$$PageEntrySize = (26 - K) + 4 = 30 - K \text{ bits} = \frac{30 - K}{8} \text{ Bytes}$$

Now we must calculate the number of entries (corresponding to the number of virtual pages) that must be kept in the page table:

$$NumberOfPages_{virtual} = \frac{AddressSpaceSize}{PageSize} = \frac{2^{32}}{2^K} = 2^{32-K}$$

The following relationship holds:

$$PageTableSize = PageEntrySize \times NumberOfPages_{virtual}$$

$$2^{17} = \frac{30 - K}{8} \times 2^{32-K}$$

$$2^{K-12} = 30 - K \Rightarrow 15 < K < 16$$

We must pick $K=16$ to stay under the specified limit of 1 Mbits of page table.

Cache, Virtual Memory and Virtual Machines

Cache and Virtual Memory



- Which address to use to index to cache sets?
 - Physical address? May be slow due to translation
 - Virtual? Protection, process switching, aliasing
- How to make a virtual cache work?
 - Page coloring
- Get the best of both: virtually indexed, physically tagged cache

Cache vs. Virtual Memory

- Concept behind VM is almost identical to concept behind cache.
- But different terminology
 - Cache: Block VM: Page
 - Cache: Cache Miss VM: Page Fault
- Caches implemented completely in hardware. VM implemented in software, with hardware support from the processor.
- Cache speeds up main memory access, while main memory speeds up VM access.

VM & Caching Comparison

Parameter	First-level cache	Virtual memory
Block (page)	16-128 bytes	4096-65,536 bytes
Hit time	1-3 clock cycles	50-150 clock cycles
Miss penalty	8-150 clock cycles	1,000,000-10,000,000 clock cycles
Access time	6-130 clock cycles	800,000-8,000,000 clock cycles
Transfer time	2-20 clock cycles	200,000-2,000,000 clock cycles
Miss rate	0.1-10%	0.00001-0.001%
Address mapping	25-45 bit physical address to 14-20 bit cache address	32-64 bit virtual address to 25-45 bit physical address



Impact of These Properties on Design

- Comparing to cache, how would we set the following design parameters?
 - block size?
 - Large, since disks perform better at transferring large blocks
 - Associativity?
 - High, to minimize miss rate
 - Write through or write back?
 - Write back, since can't afford to perform small writes to disk
- What would the impact of these choices be on:
 - miss rate
 - Extremely low. $\ll 1\%$
 - hit time
 - Must match cache/DRAM performance
 - miss latency (penalty)
 - Very high. $\sim 20\text{ms}$

Associativity of VM

- Cache miss penalty: 8-150 clock cycles
- VM miss penalty: 1,000,000 - 10,000,000 clock cycles
- Because of the high miss penalty, VM design minimizes miss rate by allowing full associativity for page placement

Write Strategies

- Disk I/O slow (millions of clock cycles)
- Always write-back; never write-through
- Use dirty bit to decide whether to write disk before eviction
- Smart disk controllers buffers writes
 - copy replaced page in buffer
 - read new page into main memory
 - write from buffer to disk

Virtual Machines

- Supports isolation and security
- Sharing a computer among many unrelated users
- Enabled by raw speed of processors, making the overhead more acceptable
- Allows different ISAs and operating systems to be presented to user programs
 - “System Virtual Machines”
 - SVM software is called “virtual machine monitor” or “hypervisor”
 - Individual virtual machines run under the monitor are called “guest VMs” or “guest OSes”

Impact of VMs on Virtual Memory

Each guest OS maintains its own set of page tables

- VMM adds a level of memory between physical and virtual memory called “real memory”
- VMM maintains shadow page table that maps guest virtual addresses to physical addresses
 - Requires VMM to detect guest’s changes to its own page table
 - Occurs naturally if accessing the page table pointer is a privileged operation