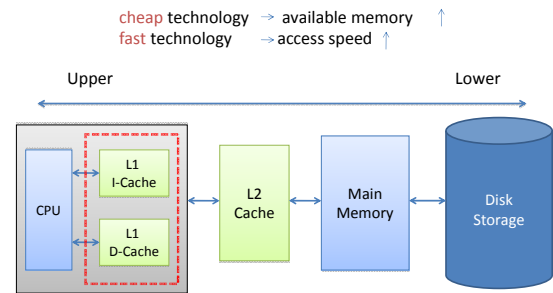


COMP4611 Tutorial 8 Memory System

1


Levels in Memory Hierarchy



Memory Hierarchy

Level	1	2	3	4
Name	Registers	Cache	Main Memory	Disk Storage
Typical size	<1KB	<16MB	<64GB	>1TB
Access time (ns)	0.25-0.5	0.5-25	50-250	5,000,000
Managed by	Compiler	Hardware	Operating System	Operating System/Operator
Speed	Fastest			Slowest
Capacity	Smallest			Biggest
Cost/bit	Highest			Lowest

iMac's PowerPC 970



Design Issues

- Variable factors affect the cache design
 - Cache size
 - Larger cache → shorter latencies
 - Cache speed and latency
 - Increasing speed shorten access latency
 - Associativity
 - one-way(direct mapping), two-way, four-way or eight-way
 - Cost
- Factors are *inter-related* → Difficult to achieve the "best cache"

Fast and large
caches are
expensive

Cache Concepts

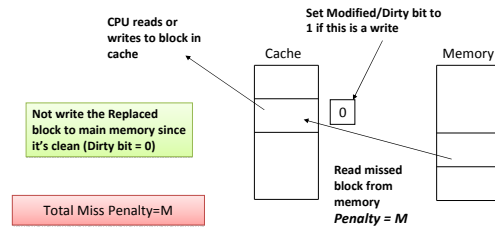
- Cache Hit
 - Data requested by the CPU is present in the cache
- Cache Miss
 - Data requested by the CPU is not present in the cache
 - On a cache miss, a block brought from the main memory may replace an existing cache block
- Hit Rate (or Hit Ratio)
 - The percentage of accesses that result in cache hits
- Cache Replacement Policies
 - FIFO (First In First Out)
 - LRU (Least Recently Used)
 - LFU (Least Frequently Used)
 - MFU (Most Frequently Used)
 - Random

Cache Concepts (cont.)

- Cache Write Policies
 - Write Through**
Data is written to both the cache block and to a block of main memory
 - Write Back**
Data is written only to the cache block ;
Modified cache block is written to main memory when it has to be replaced
- Cache Write Miss Policies
 - Write Allocate**
The cache block is allocated on a write miss, followed by write hit actions
 - No Write Allocate**
Write misses do not affect the cache;
The block is only modified in the main memory


Cache Miss Operation

- Assume: 1. Write-Back Cache with Write-Allocate
2. Block to be replaced is clean



Cache Miss Operation

- Assume: 1. Write-Back Cache with Write-Allocate
2. Block to be replaced is dirty



Example 1

Suppose a computer's address size is 64 bits (using byte addressing), the cache size is 64 Kbytes ($1\text{ K} = 2^{10}$ bytes), the block size is 64 bytes and the cache is 8-way set-associative.

Compute the following quantities:

- the number of sets in the cache
- the number of index bits
- the number of tag address bits in a block

Example 1 - Solution

- This is an 8-way set associative cache, the size of each set is $8 * \text{Block_size} = 512$ bytes
Thus, # Sets = $\text{Cache_size} / \text{Set_size} = 64\text{KB} / 512\text{ B} = 128$ Sets
- The number of index bits is determined by the number of sets.
Since there are 128 sets, 7 bits are needed as the index bits ($2^7 = 128$).
- The number of tag address bits is determined by the total address size, the number of index bits and the number of offset bits.
(Tag address bits) = # (Address Bits) - # (index bits) - # (offset bits)
(offset bits) = 6 (block size is 64 bytes and $2^6 = 64$)
Thus, # (Tag address bits) = $64 - 7 - 6 = 51$

Average Memory Access Time (AMAT)

AMAT can be expressed by Hit time, Miss rate and Miss penalty on different cache levels.

For example,

$$\text{AMAT} = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty} \quad (1\text{-level})$$

$$\text{AMAT} = \text{Hit time}_{L1} + \text{Miss rate}_{L1} \times (\text{Hit time}_{L2} + \text{Miss rate}_{L2} \times \text{Miss penalty}_{L2}) \quad (2\text{-level})$$

FYI: Not all the cases are included. e.g., 3-level cache

Example 2

- Suppose that in 1000 memory references there are 40 misses in the first level cache and 20 misses in the second level cache. What are the various miss rate?
- Assume the miss penalty from the L2 cache to memory is **100 clock cycles**, the hit time of the L2 cache is **10 clock cycles**, the hit time of L1 is **1 clock cycle**, and there are **1.5 memory references** per instruction. What is the average memory access time and average stall cycles per instruction? Ignore the impact of writes.

Example 2 - Solution

- Miss rate for L1 (either Local or Global) = $40/1000 = 4\%$

Local miss rate for L2 = $20/40 = 50\%$

Global miss rate for L2 = $20/1000 = 2\%$

- $AMAT = \text{Hit time}_{L1} + \text{Miss rate}_{L1} \times (\text{Hit time}_{L2} + \text{Miss rate}_{L2} \times \text{Miss penalty}_{L2})$
 $= 1 + 4\% \times (10 + 50\% \times 100)$
 $= 3.4 \text{ clock cycles}$

Average memory stalls per instruction

$= (AMAT - 1) \times 1.5$

$= (3.4 - 1) \times 1.5$

$= 3.6$

AMAT = 1 + Stall Cycles Per Memory Access

Virtual Memory

- Definition
 - It gives an application program the impression that it has contiguous working memory, while in fact it may be physically fragmented and may even overflow on to disk storage
- Motivations
 - Allows the same program to run in any location in physical memory
- Different terminology comparing with cache
 - Virtual memory block is called a **page**
 - A virtual memory miss is called a **page fault**

Virtual & physical address

- Physical address
 - Instruction or data address in **main memory**
 - 256 MB main memory \rightarrow 28-bit physical address
- Virtual address
 - Decided by ISA
- Virtual address = **virtual page #** and **page offset**
 - Page # identifies a particular page
 - Page offset identifies a byte within that page
- Physical address = **physical page #** and **page offset**
- Address translation
 - Virtual address issued by the processor needs to be translated into the physical address

Example 3 & Solution

The page size on a byte-addressed machine is 16 KB. The machine has 1 GB of main memory. The virtual address of the machine has 32 bits. What are the sizes of the virtual page #, physical page #, and page offset fields?

Main memory size = 2^{30} bytes,
so physical address is 30 bits.

Page size = 2^{14} bytes,
so page offset is 14 bits.

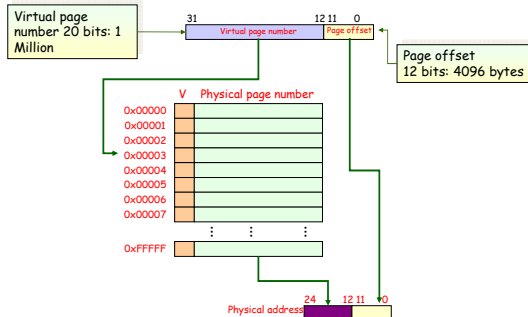
Virtual page # size = $32 - 14 = 18$ bits.

Physical page # size = $30 - 14 = 16$ bits.

Paging

- Each process has its **own page table**
- Use page number as an **index** into the page table
- Each page table entry contains the **physical page number** of the corresponding page in main memory
- A **valid bit** to indicate whether the page is in main memory or not
- A **modify bit** to indicate whether the page has been altered or not
 - If no change, the page does not have to be written to the disk when it needs to be swapped out
- Replacement policies

Page Table



TLB

- **TLB - translation lookaside buffer**

- Improve the speed of virtual address translation

Size: 8 - 4,096 entries
Hit time: 0.5 - 1 clock cycle
Miss penalty: 10 - 30 clock cycles
Miss rate: 0.01% - 1%

- A TLB entry is like a cache entry where the tag holds portions of the **virtual address** and the data portion holds the **physical page number**, protection field, **valid bit**, and **dirty bit**.

TLB (cont.)

- If the requested address is present in the TLB, the physical address can be used to access memory.
- If the requested address is not in the TLB, the translation proceeds using the page table, which is slower to access.

Overall operation of memory hierarchy

